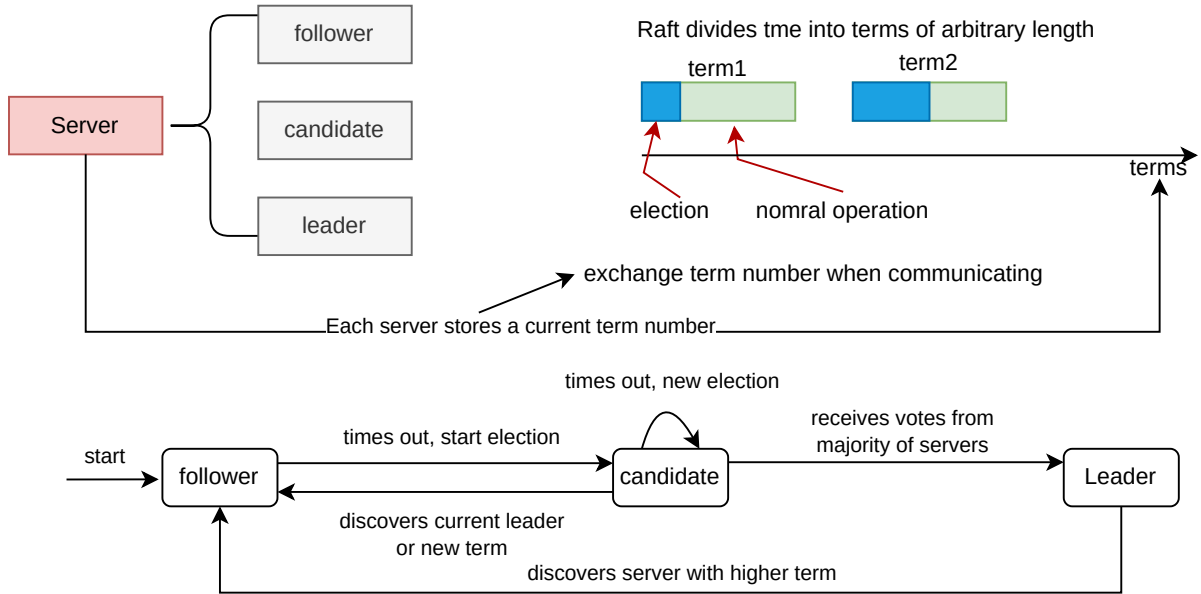
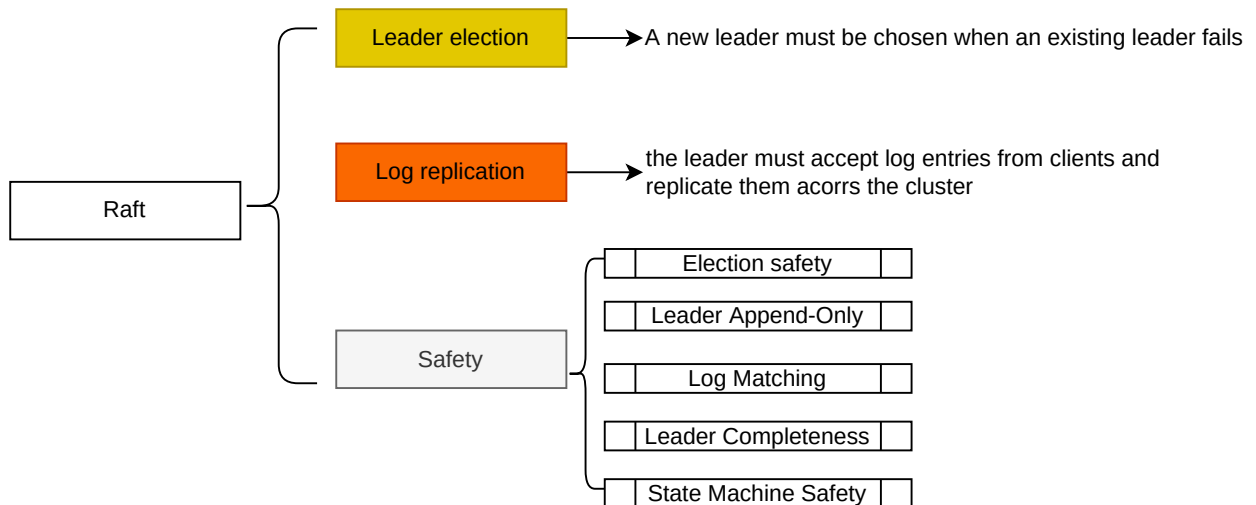


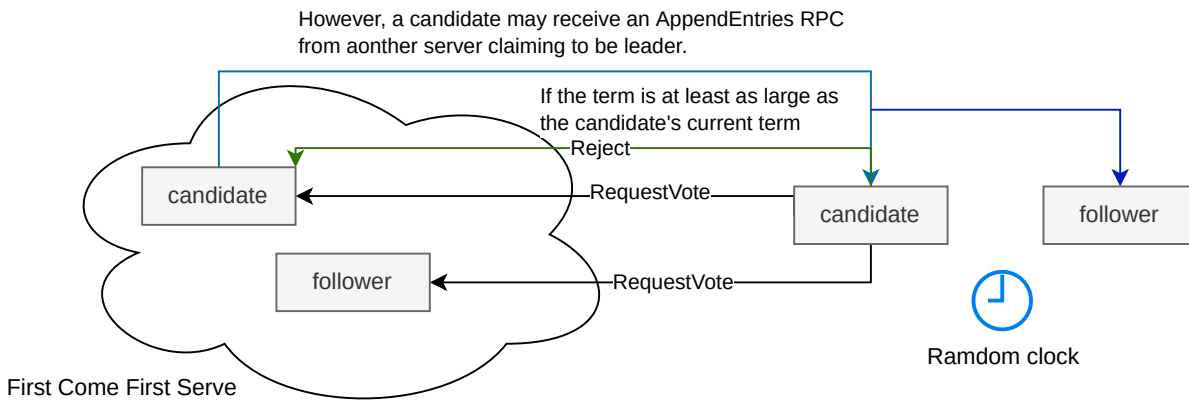
Basic



Raft Outline



The process
A follower increments its current term and transitions to candidate state.
Votes for itself and issue RequestVote RPCs in parallel to other servers
Until (a) it wins (b) timeout, no winner



A candidate must contact a majority of the cluster in order to be elected, which means that every committed entry must be present in at least one of those servers.

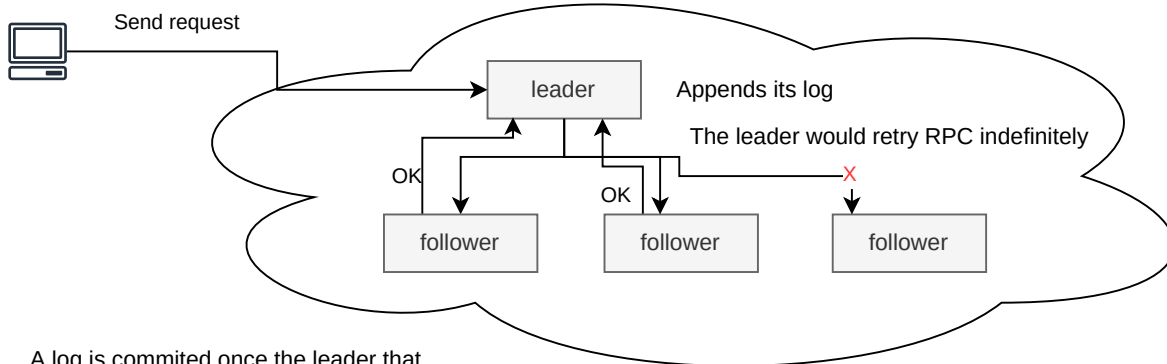


The criterion of the more up-to-date


- {

Compare the terms, greater means more up-to-date → **Election Safety**

If the terms are equal, longer log means more up-to-date



A log is committed once the leader that created the entry has replicated it on a majority of the servers.



Log Matching Property

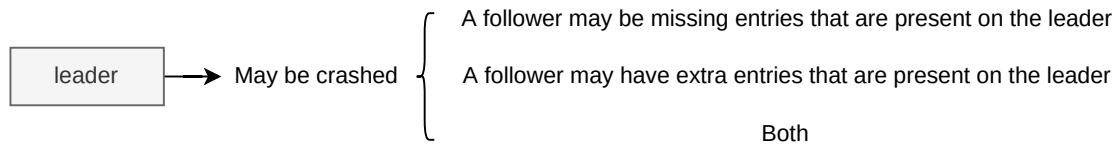
If two entries in different logs have the same index and term, then they store the same command

If two entries in different logs have the same index and term, then the logs are identical in all preceding entries.

Initial, the leader send the first log, which would must match log matching property

Induction Proof

Next, when sending the next log, leader sends the previous log's log index and term index, followers inspect it, if it finds an entry, accept it. Otherwise, refuse it



How to make the followers' log consistent with the leader

Forcing the followers' logs to duplicate the log of the leader

